



Institute of Information Technology
University of Dhaka

SE801 • SOFTWARE PROJECT LAB 3 • FINAL DEFENSE


Terra-Audit

AI-Assisted Digital MRV Platform for AWD Rice Irrigation
and Agricultural Land Management Carbon Credits

PRESENTED BY
Kazi Nahid
BSSE Roll 1437

SUPERVISED BY
Dr. Emon Kumar Dey
Professor, IIT, DU

DATE
October 2026
Final defense



2 AWD drydowns
→ 2.12 tCO₂e / ha

Verra VM0051 • VM0042

Flooded rice emits methane — and AWD reductions are hard to prove

METHANE EMISSION FACTOR

1.00 → 0.55

Alternate Wetting & Drying (AWD) —
draining the field at least twice per season
— cuts the daily methane factor by 45%
versus continuous flooding.

IPCC 2019 Table 5.12, used by Verra VM0051



Practice claims are hard to verify

Farmer irrigation logs are incomplete; sensors on every smallholder plot are too costly.



Evidence gets separated from the number

Spreadsheets lose the factors, methodology version and history behind a result.




Verification needs traceability

Verifiers and buyers must trace each tonne of CO₂e back to its evidence.

Terra-Audit links every carbon number to its evidence

An evidence-first digital MRV platform — from field boundary to reviewed export.



Rice — Alternate Wetting & Drying

Verra VM0051 v1.1

Satellite-detected drydowns cut methane; 15% QA3 uncertainty and the N₂O penalty are deducted.




Agricultural Land Management

Verra VM0042 v2.2 + VMD0054 v1.1

Measured soil-carbon change plus N₂O, fire and fuel terms, minus leakage and a non-permanence buffer; missing leakage evidence blocks the result.

Sentinel-1 radar detects AWD drydowns without field sensors



Flooded

VV z-score below -0.8

Drydown

VV rises $> 1.2\sigma$ right after a flooded date

Season

VH minimum = sowing; sharpest VH drop after the peak = harvest

Result fed to the calculation

2 AWD events • 84-day season

Live run on the deployed app • Sentinel-1 via Google Earth Engine • 12 observations, Jan–May 2026 • rule engine: AdaptiveAWDGate

Machine learning cross-checks the detector — but is not ground truth

Random Forest

	predicted		
	dry	flooded	drydown
dry	19	2	1
flooded	2	8	1
drydown	2	1	0

75.0%

agreement

0.524

macro-F1

XGBoost

	predicted		
	dry	flooded	drydown
dry	19	3	0
flooded	4	6	1
drydown	2	1	0

69.4%

agreement

0.460

macro-F1

Why add ML?

Rules are explainable but brittle. Models trained on the same Sentinel-1 features flag dates where the rules may be wrong.

Honest limitation

Models learn the gate's own labels, so this is agreement — not field accuracy. ML stays advisory until independent ground truth exists.

Rows = threshold-gate label · 3-fold stratified cross-validation

AI Validation page on the deployed app · features include vv_diff, vv_smoothed, vh, cross-ratio, days since sowing

A deterministic VM0051 engine turns evidence into tCO₂e

$$E = EF_c \times SF_w \times SF_p \times SF_o \times \text{days} \times \text{area}$$

$$\Delta CH_4 = E_{\text{baseline}} - E_{\text{project}}$$

$$\Delta CO_2e = \Delta CH_4 \times 28 \div 1000$$

$$ER = \Delta CO_2e \times (1 - 15\%) - PE_{N_2O}$$

VM0051 Eq. 6–8, 29 (QA3) · $EF_c = 0.85 \text{ kg CH}_4/\text{ha}/\text{day}$ (IPCC 2019, South Asia) · $GWP_{CH_4} = 28$ (AR5)

Worked example

1 ha · 84 days · 2 drydowns · 100 kg N/ha

Baseline CH ₄	205.5 kg
Project CH ₄	113.0 kg
Avoided CH ₄	92.5 kg
As CO ₂ e	2.59 t
– 15% uncertainty	– 0.39 t
– N ₂ O penalty	– 0.08 t

Net reduction

2.12 tCO₂e

Above 60,000 tCO₂e/yr the engine blocks instead of guessing.

WATER-REGIME FACTOR SF_w — WHERE AWD'S BENEFIT COMES FROM

1.00

Continuous flooding
0 drydowns

0.71

Single aeration
1 drydown

0.55

True AWD
≥ 2 drydowns

Worked example computed by `CarbonAssetEngine.calculate_credits()` (`src/carbon/rice.py`)

Readiness gates never let an unsupported requirement pass silently

☑️ Readiness checklist — live run, VM0051 field

satisfied

`common.monitoring_period_coverage`
Crop seasons cover the monitoring period

needs review

`common.methodology_applicability`
Declared crop outside usual scope — reviewer confirms

needs review

`common.evidence_review_status`
Field observations not yet accepted by a second person

missing

`vm0051.required_measurement_inputs`
AWD events and season length not yet supplied

unsupported

`vm0051.leakage_assessment`
Not implemented — stays blocking

unsupported

`common.additionality`
Expert determination — never automated

1 satisfied **3** needs review **1** missing **4** unsupported

Live readiness run on the deployed app (6 of 9 checks shown) · internal approval is not VVB verification or Verra issuance

AT COMMIT, THE SYSTEM FREEZES

- 🔒 Field boundary & area
- 🔒 Crop seasons & engine inputs
- 🔒 Methodology bundle & factors
- 🔒 Readiness result


INTERNAL REVIEW

submitted**in_review****internally_approved**


or changes requested / rejected / withdrawn

AI explains, code calculates, humans approve

AUTHORITATIVE • DETERMINISTIC



ADVISORY • AI



 **Bounded context**
AI only sees the active methodology bundle, evidence and calculation.

 **Validated output**
Citations must exist and numbers must match; “certified” or “approved” claims are rejected.


 **Fail closed**
One retry, then fail — nothing partial is saved.

 **Human reviewer**


No AI component can change a calculated value, mark a requirement compliant, or create registry credits

Layered components, deployed as web, API and worker


COMPONENT LAYERS



DEPLOYMENT



Solid = verified in the recorded run · dashed = configured external service · * supported hosted option



Deployment as recorded in the final report (\$5.5) · production worker and storage provenance not yet verified end to end

Design classes, persistent data and behaviour

KEY DESIGN CLASSES & MODULES

AdaptiveAWDGate

analyze_irrigation_behavior()
extract_phenology()

CarbonAssetEngine

calculate_credits()
VM0051 rice AWD

methodology.readiness

build_readiness_checklist()
per-requirement states

WorkerContext · jobs

create_job()
claim_next_job()

SpatialDataEngine

extract_clean_timeseries()
Sentinel-1 via Earth Engine

AlmCarbonEngine

calculate_credits()
VM0042 + VMD0054 leakage

carbon.snapshot

build_snapshot()
calculate_from_snapshot()

AttachmentStorage

LocalFilesystemStorage
S3Storage (strategy)

PERSISTENT DATA (50+ TABLES)



Identity

organizations · users ·
project_members



Fields & evidence

fields · crop_seasons ·
practice_events · attachments



Methodology

methodology_bundles ·
methodology_requirements ·
readiness_determinations



Calculation & review

calculations ·
review_submissions · findings



Jobs & AI


background_jobs · workers ·
ai_records

BEHAVIOUR (STATE MACHINES)

Calculation version




decided by blocking readiness items



new evidence → new version

Background job




or error (retry / invalid) · cancel_requested →
cancelled · heartbeat-monitored workers

Names as in the source code (src/, backend/worker.py) · review states shown on the readiness slide · full ER model in the report (§4)

Interface design: every user action has a visible state

Register a field



OBJECTS

map canvas • input-mode tabs


ACTION

draw / upload GeoJSON-KML / paste

STATE

boundary → confirm details → field saved

Signal analytics



OBJECTS

analysis window • detector • Run


EVENT

close < open → inline error

STATE

job pending → results + audit trail

Calculation & readiness



OBJECTS

seasons • period • Check readiness

EVENT

Check readiness → checklist

STATE

Commit → draft / ready_for_review

Shared labelled controls • inline validation • text + colour status cues • DD-MM-YYYY dates • keyboard-aware controls • responsive layouts

Code: the AWD rule engine produces an auditable trail

Date	VV z-score	Flooded	Drydown
09 Jan	1.023		
21 Jan	0.441		
02 Feb	-0.140		
14 Feb	-0.942	✓	
26 Feb	-1.440	✓	
10 Mar	-0.111		✓
22 Mar	0.390		
03 Apr	-0.606		
15 Apr	-1.326	✓	
27 Apr	-0.083		✓
09 May	0.857		
21 May	1.936		

src/signals/threshold_gate.py · AdaptiveAWDGate (excerpt)

```
# Field-level statistical baseline
vv_mean = df["vv_smoothed"].mean()
vv_std = df["vv_smoothed"].std()
if vv_std == 0 or np.isnan(vv_std):
    vv_std = 0.1 # Prevent division by zero on flat test data

# Z-score flood state detection
df["vv_zscore"] = (df["vv_smoothed"] - vv_mean) / vv_std
df["is_flooded"] = np.where(
    df["vv_zscore"] < self.z_flood_threshold, 1, 0)

# Drydown event = sharp positive jump following a flooded state
df["vv_diff"] = df["vv_smoothed"].diff().fillna(0)
threshold_jump = self.dynamic_delta_sigma * vv_std
df["drydown_event"] = np.where(
    (df["vv_diff"] > threshold_jump)
    & (df["is_flooded"].shift(1) == 1), 1, 0)
```

Left: the columns this code produces, from the live run (audit-trail table in the app) · $z_flood_threshold = -0.8$, $dynamic_delta_sigma = 1.2$

Code: the credit engine blocks instead of guessing — and a test proves it

src/carbon/rice.py · calculate_credits() (condensed)

```
# Baseline vs project CH4 (Eq. 6/8)
e_baseline = self.ef_c * 1.0 * sc_p * sc_o_bsl * days * area_ha
e_project  = self.ef_c * sf_w_project * sc_p * sc_o_wp * days * area_ha
delta_e_co2e = (e_baseline - e_project) * self.gwp_ch4 / 1000.0

# QA3 project-size gate (§8.6.3)
if delta_e_co2e > self.QA3_PROJECT_SIZE_GATE_TC02E_YR:
    return {"qa3_pathway_valid": False,
            "final_issuance": None, ...}

# 15% uncertainty, N2O penalty (Eq. 25), net (Eq. 29)
ch4_after_unc = delta_e_co2e * (1 - self.UNC_QA3_DEFAULT)
pe_n2o = (self.n2o_irrigation_penalty(q_n, area_ha)
          if awd_events > 0 else 0.0)
final_issuance = max(0.0, ch4_after_unc - pe_n2o)
```


tests/test_carbon_calculator.py

```
def test_qa3_project_size_gate_blocks_oversized_projects(
    engine):
    r = engine.calculate_credits(
        awd_events=2, season_length_days=120,
        area_ha=500_000.0, q_n_kg_per_ha=0.0)
    assert r["qa3_pathway_valid"] is False
    assert r["final_issuance"] is None
```

- Every factor is a named constant cited to its VM0051 / IPCC table
- All intermediate values are returned for the audit trail
- Snapshot calculation never re-reads the live database

Engine code condensed for the slide — full function in src/carbon/rice.py · test passes in the 177-test backend suite

177 backend tests pass — layered from equations to tenant isolation



Pass criterion

0 failures, no unjustified skips — and unsupported states must stay explicit, never silently satisfied.

Risks → contingency

No ground truth → ML advisory · unimplemented rules → blocking checks · live services → isolated fixtures.

All five checks re-run on the current code, 7–8 Oct 2026 · ESLint: 0 errors, 5 warnings

Key test cases prove gates, isolation and reproducibility

ID	Scenario	Expected	Observed	Result
TC-01	QA3 size gate — 500,000 ha field	Reject above 60,000 tCO ₂ e/yr	qa3_pathway_valid = False, no issuance	PASS
TC-03	Another organization reads a field	Cannot list or open it	empty list · direct GET → 404	PASS
TC-04	PDF after live evidence is edited	Built from the frozen snapshot	owner 200 · viewer 403 · other org 404	PASS
TC-06	AI output claims “certified”	Reject; one retry, then fail	validation error · nothing saved	PASS

Evidence: tests/test_carbon_calculator.py · tests/backend/test_fields_endpoints.py · test_evidence_verification.py · tests/test_ai_validate.py

Challenges solved — and what remains

CHALLENGE → SOLUTION



Methodology fidelity

Factors cited to tables + golden-value tests — caught a wrong EF (1.4 → 0.85).



Slow satellite & AI work

Durable job queue with a separate Python worker.



Numbers drifting after edits

Immutable snapshots; edits create a new version.



STILL OPEN

- Additionality unsupported (blocking)
- No independent AWD ground truth
- AI evaluation set incomplete
- Production provenance unverified



NEXT

- Field ground-truth campaign
- Eq. 38 uncertainty for large projects
- Remaining VM0051 terms



Thank you

Questions are welcome

Live app terra-audit.vercel.app

Repository github.com/knah1d/terra-audit

Presenter Kazi Nahid · BSSE 1437

Supervisor Dr. Emon Kumar Dey, Professor, IIT, DU



Live app



GitHub